

Deep Reinforcement Learning Algorithms

ANDREAS HOFMANN

This provides a summary and references for current standard *Deep Reinforcement Learning* (DRL) algorithms. The emphasis is on DRL for control of plants (agents) with continuous control actions, particularly, robots. As this field becomes more mature, it is important to be aware of open source tools that simplify implementation and use of the algorithms. Therefore, in addition to the algorithm descriptions, this document includes summaries of important tools, including simulators, deep learning platforms like PyTorch, and a variety of wrappers for platforms like PyTorch that make implementation easier.

1 INTRODUCTION

To begin, it is useful to understand the place of DRL within the general field of Deep Learning. The field of Deep Learning can be divided into three categories: supervised learning, unsupervised learning, and reinforcement learning.

Supervised learning requires first generating a possibly large *training set* of example inputs with corresponding outputs. The outputs are sometimes called “labels”. For example, in an image classification problem, the inputs are images, and the outputs are associated labels that describe the images. The generation of such training sets is typically the most time consuming aspect of supervised learning.

Unsupervised learning avoids the need for “supervision” by not requiring labels. It uses input (a set of images, for example), and tries to find hidden structure and relations. A typical task is to cluster inputs into sets based on some aspect of similarity.

Reinforcement learning can be considered to be “between” supervised and unsupervised learning in terms of the human supervision that it requires. Although reinforcement learning does not require direct supervision in the form of labels, it does require an environment for the *agent* to explore, and a reward function based on the agent’s behavior. The environment is often simulated. Fidelity and speed of the simulation are key determinants of success.

Intro sub sections

Overview of other sections

Use book as guide for how to organize, but bias towards continuous action spaces.

1.1 Algorithm Categorization

DRL (Deep Reinforcement Learning) Algorithms can be categorized according to a number of characteristics. These include whether the algorithm

- uses a model (*model-based* vs. *model-free*);
- learns a value function or learns a policy directly;
- is an *on-policy* or *off-policy* algorithm.

Author’s address: Andreas Hofmann, hofma@csail.mit.edu.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

© 2021 Association for Computing Machinery.

XXXX-XXXX/2021/12-ART \$15.00

<https://doi.org/10.1145/nnnnnnnn.nnnnnnnn>

With model-based RL, a model of state transition (next state given current state and action) is used. The model supports searching multiple steps ahead in order to select the optimal action, based on reward at each step. This search can be informed by both value and policy heuristics, as was done in Alpha Go [cite]. With model-free RL, no explicit state transition model is used. Instead, one or both of two simpler functions (value function and/or policy function) is learned directly.

A value function returns a scalar value for a given current state. Related to this, a Q function (or value-action function) returns a scalar value given a current state and a proposed action from this state. These scalar values represent the discounted reward, or future utility, that is expected to accrue, if a particular control policy is followed. The Q function can be used to pick the best action given a current state (the action that maximizes reward). A policy function directly returns an action for a given current state. Ideally, the policy function is learned well enough that it closely approximates the *optimal* policy function (returns the best action for a given current state).

An *actor-critic* approach uses both value and policy functions.

In order to understand the difference between on-policy and off-policy methods, it is important to consider and separate out the policy being learned and the policy used to generate the data. For on-policy methods, these are the same. For off-policy methods, these are not the same. For example, in Q-learning (value function and Q function learning), the policy used to control the agent in the environment (to generate the data) is only sometimes the policy being learned. In particular, the action is sometimes chosen from the policy being learned (greedy option), and is sometimes chosen randomly (exploratory option). This helps to ensure that the learning does not prematurely converge to a local optimum that is significantly less optimal than the global optimum. Selection of whether to choose the greedy or exploratory action is based on an *epsilon* factor that decays over time, so that as the policy is learned, selection of the greedy action becomes more likely (leveraging the fact that the policy should be getting better over time).

Q-Learning is an example of an off-policy method. SARSA is an example of an on-policy method. See also <https://analyticsindiamag.com/reinforcement-learning-policy>.

See also <https://smartlabai.medium.com/reinforcement-learning-algorithms-an-intuitive-overview-904e2dff5bbc>

<https://www.guru99.com/reinforcement-learning-tutorial.html>

https://en.wikipedia.org/wiki/Deep_reinforcement_learning#Off-policy_reinforcement_learning

<https://analyticsindiamag.com/reinforcement-learning-policy>.

Additional characteristics used to categorize DRL algorithms include:

- whether the output action is discrete or continuous;
- whether the input is raw pixels or a state estimate vector;
- whether the system is deterministic or stochastic.

[Investigate whether there are algorithms that have hybrid discrete/continuous output actions.]

In principle, DRL algorithms model Markov Decision Processes (MDPs). As such, they are based on the notion of state. However, in many cases, the distinction between state and sensor data is blurred. For example, in many DRL algorithms, an environment like the cart-pole environment in the Open AI gym provides pixels as the state, rather than the traditional position and angle state estimate, even though, from a traditional control systems standpoint, pixels should be regarded as sensor input. This blurring can be considered a feature as it allows for end to end control algorithms.

Regarding whether the system is deterministic or stochastic, a general MDP supports a stochastic environment, where a particular control action given a current state implies a distribution of possible next states. The deterministic case is just the degenerate case of the MDP, where the control action implies a single next state. There is some confusion about this terminology in some of the DRL algorithms; in some cases, stochastic implies

a control policy that provides a distribution of possible actions given the current state. Thus, it is possible to have both a distribution of possible actions, and a distribution of next states conditioned for each action.

1.2 Algorithm Summaries

Intro to algorithms (summary of each) List algorithms, organized (see web sites)
 deterministic vs. stochastic, MDP
 Confusion about this also

1.3 To Do

Focus on continuous action models like mountain car.

Think about whether it makes sense to have hybrid discrete/continuous actions, maybe from two separate heads of the DNN.

2 BACKGROUND

Problem statement, take from below.

3 DQN

Q value algorithm

3.1 Agent, Reward, Value Function, Q Function, and Policy

In Reinforcement Learning, an *agent* takes actions in an environment in order to accumulate *rewards*. The immediate reward at time increment t is given as R_t . When learning how an agent should behave, it is useful to consider *cumulative* reward, rather than just immediate reward. Cumulative reward is defined as the total reward accumulated by the agent over some time horizon:

$$C_t = R_t + R_{t+1} + R_{t+2} + \dots = \sum_{k=0}^T R_{t+k} \quad (1)$$

Here, the agent takes actions at regular time increments, denoted by k . The time horizon, T , is specified in terms of the maximum value for k .

In practice, when learning agent behaviors, it is useful to favor immediate reward over future reward. This is expressed as *discounted* reward, using a discount rate, γ :

$$G_t = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \dots = \sum_{k=0}^T \gamma^k R_{t+k} \quad (2)$$

With the discount rate, T is usually set at infinity

$$G_t = \sum_{k=0}^{\infty} \gamma^k R_{t+k} \quad (3)$$

The discount rate implies a time horizon in that rewards diminish significantly in the distant future.

A policy, π , specifies an action, a , to be taken by the agent when the state is s . Two important functions can be defined with respect to a policy. First, a value function gives the total expected reward given a current state, s_t , if a particular policy, π , is followed.

$$v_{\pi}(s) = \mathbb{E}(G_t | S_t = s) \quad (4)$$

Second, an action-value or Q function gives the total expected reward given a current state, s_t , a current action, A_t , and a particular policy, π , to be followed (for future actions).

$$q_\pi(s, a) = \mathbb{E}(G_t | S_t = s, A_t = a) \quad (5)$$

The policy can be expressed in terms of the Q function as

$$\pi(s) = \operatorname{argmax}_a q_\pi(s, a) \quad (6)$$

Note that this is circular; the Q function is, itself, a function of the policy. Thus, the challenge is how to compute (learn) the Q function. To accomplish this, we leverage the Bellman equation:

$$q_\pi(s, a) = r + \gamma q_\pi(s', \pi(s')) = r + \gamma \operatorname{argmax}_{a'} q_\pi(s', a') \quad (7)$$

where r is the immediate reward for action a , s' is the next state resulting from taking action a at state s , and a' is the action to be taken at state s' .

3.2 Training: Loss Function, Agent Steps, Replay Memory, and Policy Update

The *temporal difference error*, δ , is defined as the difference between the two sides of 8:

$$\delta = q_\pi(s, a) - \left(r + \gamma \operatorname{argmax}_{a'} q_\pi(s', a') \right) \quad (8)$$

During training, RL algorithms strive to update $q_\pi(s, a)$ so that δ approaches 0. In (non-deep) RL, $q_\pi(s, a)$ is represented as a table, and is updated by:

$$q_\pi(s, a) = q_\pi(s, a) + \alpha \left(r + \gamma \operatorname{argmax}_{a'} q_\pi(s', a') - q_\pi(s, a) \right) \quad (9)$$

where α is the learning rate.

In Deep Reinforcement Learning, $q_\pi(s, a)$ is represented using a Deep Neural Net rather than a table. Training (minimization of δ) is accomplished by defining a loss based on δ . In the PyTorch DQN tutorial (https://pytorch.org/tutorials/intermediate/reinforcement_q_learning.html), a *Huber Loss* function is used. This loss, \mathcal{L} , is defined as

$$\mathcal{L} = \begin{cases} \frac{1}{2} \delta^2 & \text{for } |\delta| \leq 1, \\ |\delta| - \frac{1}{2} & \text{otherwise.} \end{cases} \quad (10)$$

Training occurs over some specified number of episodes (50, for example). Within each episode, the environment state is first re-initialized, and then the agent explores the environment by taking *steps*. For each step, the agent takes an action, resulting in a reward, and a transition from the current state to the next state. Thus, a step can be summarized by the tuple $\langle \text{state action next-state reward} \rangle$. An episode ends when the environment decides this (usually because the agent has entered a failure state), or at some specified limit on the number of steps per episode.

After each step, the result tuple is entered into *replay memory*. Replay memory is a buffer of these tuples with a limited size (10000 for example). As result tuples are added to the end of replay memory, old results are dropped from the beginning. Thus, replay memory stores the last n (10000 for example) results of experiment steps taken by the agent.

After each agent step, the Q model is updated (optimized). This is accomplished not by using the results from the most recent step, but rather, by sampling a *batch* of results from replay memory. A typical batch size is 128. Leveraging the Huber loss function defined above, the loss for the entire batch is defined as

$$\mathcal{L}_B = \frac{1}{|B|} \sum_{(s,a,s',r) \in B} \mathcal{L} \quad (11)$$

Model weights are then updated using standard back propagation, based on this loss for the batch.

Experience replay memory ensures that the experience samples in a batch are not correlated; that they are not overly biased towards results from the most recent policy (model weights). This stabilizes and improves DQN training.

As in traditional reinforcement learning, the action selected by an agent during a step is sometimes *greedy* (following the current learned policy), and is sometimes *exploratory* (random). The probability of selecting a random action starts high (at 0.9, for example), and then decays exponentially to a small value (0.05, for example).

3.3 Policy and Target Networks

During training, two parallel networks are used: *policy* network, and *target* network. The policy network weights are updated at each training step, using standard back propagation. The target network weights are kept fixed for most steps; they are updated at a specified interval (every 10 training steps for example) from the policy network weights. The policy network is used to compute the left side of 8. The target network is used to compute $q_\pi(s', a')$ for the right side of 8. Using a separate target network in this way improves training stability.

3.4 Template Styles

`\documentclass[STYLE]{acmart}`

Journals use one of three template styles. All but three ACM journals use the `acmsmall` template style:

- `acmsmall`: The default journal template style.
- `acmlarge`: Used by JOCCH and TAP.
- `acmtog`: Used by TOG.

4 TABLES

The “`acmart`” document class includes the “`booktabs`” package — <https://ctan.org/pkg/booktabs> — for preparing high-quality tables.

Table captions are placed *above* the table.

Because tables cannot be split across pages, the best placement for them is typically the top of the page nearest their initial cite. To ensure this proper “floating” placement of tables, use the environment **`table`** to enclose the table’s contents and the table caption. The contents of the table itself must go in the **`tabular`** environment, to be aligned properly in rows and columns, with the desired horizontal and vertical rules. Again, detailed instructions on **`tabular`** material are found in the *L^AT_EX User’s Guide*.

Immediately following this sentence is the point at which Table 1 is included in the input file; compare the placement of the table here with the table in the printed output of this document.

To set a wider table, which takes up the whole width of the page’s live area, use the environment **`table*`** to enclose the table’s contents and the table caption. As with a single-column table, this wide table will “float” to a location deemed more desirable. Immediately following this sentence is the point at which Table 2 is included in the input file; again, it is instructive to compare the placement of the table here with the table in the printed output of this document.

Table 1. Frequency of Special Characters

Non-English or Math	Frequency	Comments
Ø	1 in 1,000	For Swedish names
π	1 in 5	Common in math
\$	4 in 5	Used in business
Ψ_1^2	1 in 40,000	Unexplained usage

Table 2. Some Typical Commands

Command	A Number	Comments
<code>\author</code>	100	Author
<code>\table</code>	300	For tables
<code>\table*</code>	400	For wider tables

5 MATH EQUATIONS

You may want to display math equations in three distinct styles: inline, numbered or non-numbered display. Each of the three are discussed in the next sections.

5.1 Inline (In-text) Equations

A formula that appears in the running text is called an inline or in-text formula. It is produced by the **math** environment, which can be invoked with the usual `\begin . . . \end` construction or with the short form `$. . . $`. You can use any of the symbols and structures, from α to ω , available in \LaTeX [20]; this section will simply show a few examples of in-text equations in context. Notice how this equation: $\lim_{n \rightarrow \infty} x = 0$, set here in in-line math style, looks slightly different when set in display style. (See next section).

5.2 Display Equations

A numbered display equation—one set off by vertical space from the text and centered horizontally—is produced by the **equation** environment. An unnumbered display equation is produced by the **displaymath** environment.

Again, in either environment, you can use any of the symbols and structures available in \LaTeX ; this section will just give a couple of examples of display equations in context. First, consider the equation, shown as an inline equation above:

$$\lim_{n \rightarrow \infty} x = 0 \quad (12)$$

Notice how it is formatted somewhat differently in the **displaymath** environment. Now, we'll enter an unnumbered equation:

$$\sum_{i=0}^{\infty} x + 1$$

and follow it with another numbered equation:

$$\sum_{i=0}^{\infty} x_i = \int_0^{\pi+2} f \quad (13)$$

just to demonstrate \LaTeX 's able handling of numbering.

6 FIGURES

The “figure” environment should be used for figures. One or more images can be placed within a figure. If your figure contains third-party material, you must clearly identify it as such, as shown in the example below.



Fig. 1. 1907 Franklin Model D roadster. Photograph by Harris & Ewing, Inc. [Public domain], via Wikimedia Commons. (<https://goo.gl/VLCRBB>).

Your figures should contain a caption which describes the figure to the reader. Figure captions go below the figure. Your figures should **also** include a description suitable for screen readers, to assist the visually-challenged to better understand your work.

Figure captions are placed *below* the figure.

7 CITATIONS AND BIBLIOGRAPHIES

The use of `BIBTEX` for the preparation and formatting of one’s references is strongly recommended. Authors’ names should be complete — use full first names (“Donald E. Knuth”) not initials (“D. E. Knuth”) — and the salient identifying features of a reference should be included: title, year, volume, number, pages, article DOI, etc.

The bibliography is included in your source document with these two commands, placed just before the `\end{document}` command:

```
\bibliographystyle{ACM-Reference-Format}
\bibliography{bibfile}
```

where “bibfile” is the name, without the “.bib” suffix, of the `BIBTEX` file.

Citations and references are numbered by default. A small number of ACM publications have citations and references formatted in the “author year” style; for these exceptions, please include this command in the **preamble** (before “`\begin{document}`”) of your `TEX` source:

```
\citestyle{acmauthoryear}
```

Some examples. A paginated journal article [2], an enumerated journal article [7], a reference to an entire issue [6], a monograph (whole book) [19], a monograph/whole book in a series (see 2a in spec. document) [13], a divisible-book such as an anthology or compilation [9] followed by the same example, however we only output the series if the volume number is given [10] (so Editor00a’s series should NOT be present since it has no vol. no.), a chapter in a divisible book [30], a chapter in a divisible book in a series [8], a multi-volume work as book [18], an article in a proceedings (of a conference, symposium, workshop for example) (paginated proceedings article) [3], a proceedings article with all possible elements [29], an example of an enumerated proceedings article [11], an informally published work [12], a doctoral dissertation [5], a master’s thesis: [4], an online document / world wide web resource [1, 24, 31], a video game (Case 1) [23] and (Case 2) [22] and [21] and (Case 3) a patent [28], work accepted for publication [25], ‘YYYYb’-test for prolific author [26] and [27]. Other cites might contain ‘duplicate’ DOI and URLs (some SIAM articles) [17]. Boris / Barbara Beeton: multi-volume works as books [15] and [14]. A couple of citations with DOIs: [16, 17]. Online citations: [31?, 32].

8 SIGCHI EXTENDED ABSTRACTS

The “sigchi-a” template style (available only in `TEX` and not in Word) produces a landscape-orientation formatted article, with a wide left margin. Three environments are available for use with the “sigchi-a” template style, and produce formatted output in the margin:

- sidebar: Place formatted text in the margin.
- marginfigure: Place a figure in the margin.
- margintable: Place a table in the margin.

ACKNOWLEDGMENTS

To Robert, for the bagels and explaining CMYK and color spaces.

REFERENCES

- [1] Rafal Ablamowicz and Bertfried Fauser. 2007. *CLIFFORD: a Maple 11 Package for Clifford Algebra Computations, version 11*. Retrieved February 28, 2008 from <http://math.tntech.edu/rafal/cliff11/index.html>
- [2] Patricia S. Abril and Robert Plant. 2007. The patent holder’s dilemma: Buy, sell, or troll? *Commun. ACM* 50, 1 (Jan. 2007), 36–44. <https://doi.org/10.1145/1188913.1188915>
- [3] Sten Andler. 1979. Predicate Path expressions. In *Proceedings of the 6th. ACM SIGACT-SIGPLAN symposium on Principles of Programming Languages (POPL ’79)*. ACM Press, New York, NY, 226–236. <https://doi.org/10.1145/567752.567774>
- [4] David A. Anisi. 2003. *Optimal Motion Control of a Ground Vehicle*. Master’s thesis. Royal Institute of Technology (KTH), Stockholm, Sweden.

- [5] Kenneth L. Clarkson. 1985. *Algorithms for Closest-Point Problems (Computational Geometry)*. Ph. D. Dissertation. Stanford University, Palo Alto, CA. UMI Order Number: AAT 8506171.
- [6] Jacques Cohen (Ed.). 1996. Special issue: Digital Libraries. *Commun. ACM* 39, 11 (Nov. 1996).
- [7] Sarah Cohen, Werner Nutt, and Yehoshua Sagie. 2007. Deciding equivalences among conjunctive aggregate queries. *J. ACM* 54, 2, Article 5 (April 2007), 50 pages. <https://doi.org/10.1145/1219092.1219093>
- [8] Bruce P. Douglass, David Harel, and Mark B. Trakhtenbrot. 1998. Statecharts in use: structured analysis and object-orientation. In *Lectures on Embedded Systems*, Grzegorz Rozenberg and Frits W. Vaandrager (Eds.). Lecture Notes in Computer Science, Vol. 1494. Springer-Verlag, London, 368–394. https://doi.org/10.1007/3-540-65193-4_29
- [9] Ian Editor (Ed.). 2007. *The title of book one* (1st. ed.). The name of the series one, Vol. 9. University of Chicago Press, Chicago. <https://doi.org/10.1007/3-540-09237-4>
- [10] Ian Editor (Ed.). 2008. *The title of book two* (2nd. ed.). University of Chicago Press, Chicago, Chapter 100. <https://doi.org/10.1007/3-540-09237-4>
- [11] Matthew Van Gundy, Davide Balzarotti, and Giovanni Vigna. 2007. Catch me, if you can: Evading network signatures with web-based polymorphic worms. In *Proceedings of the first USENIX workshop on Offensive Technologies (WOOT '07)*. USENIX Association, Berkley, CA, Article 7, 9 pages.
- [12] David Harel. 1978. *LOGICS of Programs: AXIOMATICS and DESCRIPTIVE POWER*. MIT Research Lab Technical Report TR-200. Massachusetts Institute of Technology, Cambridge, MA.
- [13] David Harel. 1979. *First-Order Dynamic Logic*. Lecture Notes in Computer Science, Vol. 68. Springer-Verlag, New York, NY. <https://doi.org/10.1007/3-540-09237-4>
- [14] Lars Hörmander. 1985. *The analysis of linear partial differential operators. III*. Grundlehren der Mathematischen Wissenschaften [Fundamental Principles of Mathematical Sciences], Vol. 275. Springer-Verlag, Berlin, Germany. viii+525 pages. Pseudodifferential operators.
- [15] Lars Hörmander. 1985. *The analysis of linear partial differential operators. IV*. Grundlehren der Mathematischen Wissenschaften [Fundamental Principles of Mathematical Sciences], Vol. 275. Springer-Verlag, Berlin, Germany. vii+352 pages. Fourier integral operators.
- [16] IEEE. 2004. IEEE TCSC Executive Committee. In *Proceedings of the IEEE International Conference on Web Services (ICWS '04)*. IEEE Computer Society, Washington, DC, USA, 21–22. <https://doi.org/10.1109/ICWS.2004.64>
- [17] Markus Kirschmer and John Voight. 2010. Algorithmic Enumeration of Ideal Classes for Quaternion Orders. *SIAM J. Comput.* 39, 5 (Jan. 2010), 1714–1747. <https://doi.org/10.1137/080734467>
- [18] Donald E. Knuth. 1997. *The Art of Computer Programming, Vol. 1: Fundamental Algorithms (3rd. ed.)*. Addison Wesley Longman Publishing Co., Inc.
- [19] David Kosiur. 2001. *Understanding Policy-Based Networking* (2nd. ed.). Wiley, New York, NY.
- [20] Leslie Lamport. 1986. *LT_EX: A Document Preparation System*. Addison-Wesley, Reading, MA.
- [21] Newton Lee. 2005. Interview with Bill Kinder: January 13, 2005. Video. *Comput. Entertain.* 3, 1, Article 4 (Jan.-March 2005). <https://doi.org/10.1145/1057270.1057278>
- [22] Dave Novak. 2003. Solder man. Video. In *ACM SIGGRAPH 2003 Video Review on Animation theater Program: Part I - Vol. 145 (July 27–27, 2003)*. ACM Press, New York, NY, 4. <https://doi.org/99.9999/woot07-S422>
- [23] Barack Obama. 2008. A more perfect union. Video. Retrieved March 21, 2008 from <http://video.google.com/videoplay?docid=6528042696351994555>
- [24] Poker-Edge.Com. 2006. Stats and Analysis. Retrieved June 7, 2006 from <http://www.poker-edge.com/stats.php>
- [25] Bernard Rous. 2008. The Enabling of Digital Libraries. *Digital Libraries* 12, 3, Article 5 (July 2008). To appear.
- [26] Mehdi Saeedi, Morteza Saheb Zamani, and Mehdi Sedighi. 2010. A library-based synthesis methodology for reversible logic. *Microelectron. J.* 41, 4 (April 2010), 185–194.
- [27] Mehdi Saeedi, Morteza Saheb Zamani, Mehdi Sedighi, and Zahra Sasanian. 2010. Synthesis of Reversible Circuit Using Cycle-Based Approach. *J. Emerg. Technol. Comput. Syst.* 6, 4 (Dec. 2010).
- [28] Joseph Scientist. 2009. The fountain of youth. Patent No. 12345, Filed July 1st., 2008, Issued Aug. 9th., 2009.
- [29] Stan W. Smith. 2010. An experiment in bibliographic mark-up: Parsing metadata for XML export. In *Proceedings of the 3rd. annual workshop on Librarians and Computers (LAC '10, Vol. 3)*, Reginald N. Smythe and Alexander Noble (Eds.). Paparazzi Press, Milan Italy, 422–431. <https://doi.org/99.9999/woot07-S422>
- [30] Asad Z. Spector. 1990. Achieving application requirements. In *Distributed Systems* (2nd. ed.), Sape Mullender (Ed.). ACM Press, New York, NY, 19–33. <https://doi.org/10.1145/90417.90738>
- [31] Harry Thornburg. 2001. *Introduction to Bayesian Statistics*. Retrieved March 2, 2005 from <http://ccrma.stanford.edu/~jos/bayes/bayes.html>
- [32] TUG. 2017. *Institutional members of the T_EX Users Group*. Retrieved May 27, 2017 from <http://wwwtug.org/instmem.html>
- [33] JCTANacmart Boris Veytsman. [n. d.]. *acmart—Class for typesetting publications of ACM*. Retrieved May 27, 2017 from <http://www.ctan.org/pkg/acmart>

A USEFUL WEB REFERENCES

A.1 Introduction to Deep Learning

<http://incompleteideas.net/book/the-book.html>

<https://markus-x-buchholz.medium.com/deep-reinforcement-learning-introduction-deep-q-network-dqn-algorithm-fb74bf4d6862>

<https://www.kdnuggets.com/2018/03/5-things-reinforcement-learning.html>

<https://www.guru99.com/reinforcement-learning-tutorial.html>

<https://deeppmind.com/blog/article/deep-reinforcement-learning>

<https://smartlabai.medium.com/reinforcement-learning-algorithms-an-intuitive-overview-904e2dff5bbc>

https://en.wikipedia.org/wiki/Deep_reinforcement_learning#Off-policy_reinforcement_learning

A.2 DQN

https://pytorch.org/tutorials/intermediate/reinforcement_q_learning.html

<https://unnatsingh.medium.com/deep-q-network-with-pytorch-d1ca6f40bfda>

<https://www.toptal.com/deep-learning/pytorch-reinforcement-learning-tutorial>

<https://towardsdatascience.com/deep-q-learning-tutorial-mindqn-2a4c855abffc>

<https://towardsdatascience.com/dqn-part-1-vanilla-deep-q-networks-6eb4a00febf8>

<https://en.wikipedia.org/wiki/Q-learning>

<https://openai.com/blog/openai-baselines-dqn/>

A.3 General Algorithms and Repositories

<https://github.com/p-christ/Deep-Reinforcement-Learning-Algorithms-with-PyTorch>

<https://github.com/openai/baselines>

A.4 Reference Papers and Surveys

<https://spinningup.openai.com/en/latest/spinningup/keypapers.html>

<https://arxiv.org/abs/1906.10025>

A.5 Latex

<https://www.latextemplates.com/template/acm-publications>

<http://math.mit.edu/~dspivak/files/symbols-all.pdf>